

Automated Stock Market Investment Using Reinforcement Learning (Part 2)

Ankan, Ishan, Bijayan, Shreyam, Praveen

May 22, 2024

Contents

- 1 Previous Works
- 2 Current Work
- 3 First Approach on DQN
- 4 Second Approach on DQN
- 5 Future Plans
- 6 Final work
- 7 Main idea of the DQN training
- 8 Final code summary
- 9 Outputs

Previous Works

- We implemented the Z-score algorithm on three stocks: Google, Tesla, and TCS.
- Following that, we implemented the Q-learning algorithm on the same set of stocks.
- Subsequently, we conducted a comparative analysis of the outcomes generated by the two aforementioned methods.

Current Work

We implemented two variations of the Deep Q-Network (DQN) method, each tailored to prioritize different types of rewards. We used the stock of Google for this.



Figure: Stock Values of Google over train and test set

First Approach on DQN

- The reward that we chose for this method is
 $reward = \max(stock_value(t) - bought_price, 0)$
- We appended the tuple $(state, action, reward, next_state)$ to a list we called memory, from where we intend to sample a batch to perform the learning of the neural network.
- The actions are chosen epsilon greedily.

First Approach on DQN




Figure: Google Stock

Second Approach on DQN

- We trained the neural network by sampling mini-batches from memory and updating the Q-values using the Bellman equation.
- We then updated the target Q-network periodically.
- Then we decreased epsilon over time to gradually shift from exploration to exploitation.
- Then we tracked the total reward and loss over a certain number of epochs.

Second Approach on DQN

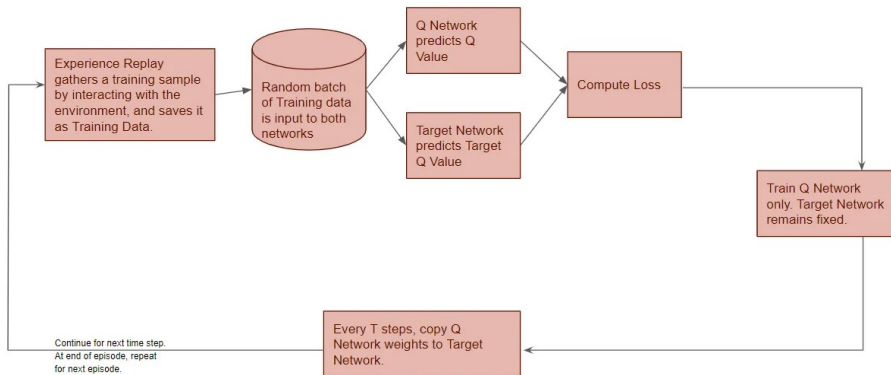


Figure: Process of the DQN

Second Approach on DQN




Figure: Total loss (left) and Total reward (right) on the stock of Google

Second Approach on DQN

DQN: train s-reward 76, profits 991, test s-reward 6, profits 389




Figure: Profit and Rewards on Stock data of Google with actions

Final work

- 1 We have extended the idea of the previous DQN algorithm for single stock.
- 2 We extended it to multi stocks case.
- 3 We have tested our code for "Google, Tesla, Apple, Microsoft, Amazon" stocks at once.
- 4 First we looked into how the untrained neural net model is working.
- 5 Then we trained the neural net model using Deep Q-learning technique to finally predict the stock market with more accuracy.

Main idea of the DQN training

- 1 We start with a two neural network one being the target.
- 2 We will run the episodes and for each step, we will concat the observations, along with the actions taken and the rewards to a memory
- 3 Once the memory has reached a certain size, we will perform a batch sampling.
- 4 Now considering the target network, we will define target values for each observation.
- 5 Feeding the states along with the target values in the current network, we will train the network.
- 6 Finally we will copy the weights of the current network to the target network and we will continue the process until done or we enter a None state.

Final code summary I

- Introduction:

- ▶ The code implements a Deep Q-Learning (DQN) model for stock trading.
- ▶ TensorFlow and Keras are used for building and training the neural network model.
- ▶ Historical stock market data is obtained using the Yahoo Finance API.

- Code Overview:

- ▶ Data Collection and Preprocessing:

- ★ Historical stock market data is collected using the `yfinance` library.
- ★ Data preprocessing involves filling missing values and normalizing the data.

- ▶ Environment Setup:

- ★ An environment class (`StockMarketEnv`) is defined to simulate the stock trading environment.
- ★ Methods for resetting, stepping through actions, and computing rewards are included.

- ▶ Deep Q-Network (DQN):

Final code summary II

- ★ The DQN model consists of a neural network architecture with three dense layers.
- ★ It utilizes a target network and experience replay buffer for stable and efficient learning.
- ▶ Training and Testing:
 - ★ The training process involves iteratively interacting with the environment and updating Q-values based on rewards.
 - ★ The trained model is tested by simulating trading on test data, and the net worth over time is plotted.
- Results:
 - ▶ Net Worth Over Time:
 - ★ The net worth of the agent is plotted over time, starting from the initial investment amount.
 - ▶ Stock Ownership Over Time:
 - ★ The ownership of each stock over time is plotted to visualize the agent's investment decisions.
 - ▶ Final Net Worth:
 - ★ The final net worth achieved by the agent after trading is displayed.

Final code summary III

- Conclusion:

- ▶ The code demonstrates the application of DQN for stock trading, showcasing the potential of RL techniques in financial decision-making.
- ▶ By leveraging historical data and reinforcement learning, the model learns to make investment decisions that optimize the agent's net worth over time.

Outputs I

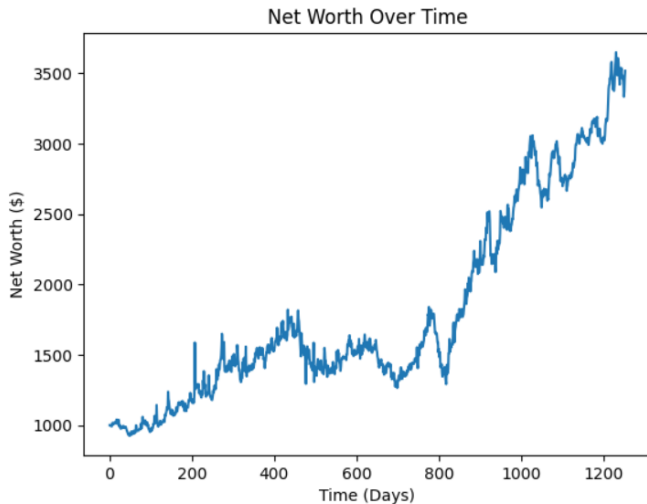
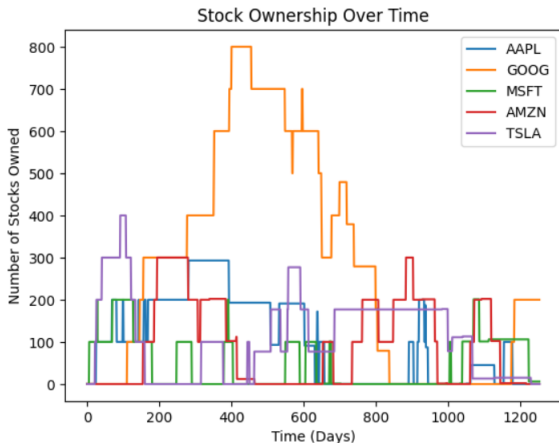


Figure: Net worth graph

Outputs II



Initial investment amount \$ 1000

Stocks investment done in: AAPL GOOG MSFT AMZN TSLA

Final Net Worth: \$3517.90

Figure: Stock Ownership graph